

MTH 9821 Numerical Methods for Finance
Fall 2023
Homework 5

1 Backtest overfitting

1. [Writing] Show that the maximum function, $\max : \mathbb{R}^n \rightarrow \mathbb{R}$, is convex.
2. [Writing] An equity analyst at a hedgefund is developing a trading strategy on a stock S with price S_t at time t .

He has a signal α_t as well, which we uses to determine how much to long/short in the next period. In other words, his holdings at time t are given by $f(\alpha_t, S_t)$, for some “trading strategy” f .

The strategy’s PnL is then given by

$$\text{PnL}(f) = \sum_{t=0}^N f(\alpha_t, S_t)(S_{t+1} - S_t)$$

where $\{\alpha_t\}_{t=0\dots N}$ and $\{S_t\}_{t=0\dots N}$ are the historical signal and stock price series and we assume interest rates are zero.

The analyst backtests $n > 1$ trading strategies $\{f_i : i = 1 \dots n\}$, and picks best performing one f_* . He tells his boss that his strategy can make $\text{PnL}(f_*)$.

Show using part (1) that *on average*, this is an overestimate.

2 Polynomial regression

[Coding] In this exercise, you will implement polynomial regression. Fill in all the missing pieces of code in the provided [regression.cpp](#).

1. [Reading] Read *Lecture 4. The Singular Value Decomposition* and *Lecture 11. Least Squares Problems* in Trefethen’s Numerical Linear Algebra [TB97] for a concise overview of SVD and polynomial regression.

2. Calculate the Vandermonde matrix given by

$$\mathcal{V}(x) = \begin{pmatrix} 1 & x_1 & x_1^2 & \dots & x_1^{\deg} \\ 1 & x_2 & x_2^2 & \dots & x_2^{\deg} \\ 1 & x_3 & x_3^2 & \dots & x_3^{\deg} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_n & x_n^2 & \dots & x_n^{\deg} \end{pmatrix}$$

3. For the least squares fit, we will solve the so-called *normal equation*

$$X'X\beta = X'y \quad (1)$$

As discussed in the lecture, the matrices $X = \mathcal{V}(x)$ have a tendency to be ill-conditioned¹.

To improve numerical stability, we *precondition* the matrix X by dividing each column by its ℓ^2 norm:

$$\hat{X} = XD^{-1}$$

where

$$D_{ij} = \delta_{ij} \sqrt{\sum_{k=1}^n x_{ki}^2}$$

Note: you should not be calculating D as a dense matrix in Eigen.

4. Use Eigen's BDCSVD solver to calculate the singular value decomposition of \hat{X} . To save time, be sure to pass the flags `ComputeThinU` | `ComputeThinV` to the solver.

$$\hat{X} = U\Sigma V'$$

The coefficients and fitted values are found by

$$\beta = D^{-1}V\Sigma^{-1}U'y \quad (2)$$

$$\hat{y} = UU'y \quad (3)$$

Note: the calculation of $U'y$ can be shared between β and \hat{y} . Don't calculate \hat{y} as $X\beta$ as this may lead to unnecessary loss of precision.

[Writing] Prove that β as given in equation (2) satisfies equation (1) and that \hat{y} as given in equation (3) equals $X\beta$. Do not assume in your proof that either U or V are square matrices or that X has full rank.

Explain (with math and formulas) the meaning of the `ComputeThinU`, `ComputeThinV` flags.

¹if the values x aren't uniformly distributed on the unit circle in \mathbb{C}

5. Calculate the pseudo Vandermonde matrix for Hermite polynomials given by

$$\mathcal{H}(x) = \begin{pmatrix} 1 & x_1 & x_1^2 - 1 & \dots & \text{He}_{\deg}(x_1) \\ 1 & x_2 & x_2^2 - 1 & \dots & \text{He}_{\deg}(x_2) \\ 1 & x_3 & x_3^2 - 1 & \dots & \text{He}_{\deg}(x_3) \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_n & x_n^2 - 1 & \dots & \text{He}_{\deg}(x_n) \end{pmatrix}$$

If x roughly follows a distribution $\mathcal{N}(\mu, \sigma^2)$, the columns of $\mathcal{H}(x)$ can be made roughly orthogonal by standardizing:

$$\hat{x}_i = \frac{x_i - \text{mean}(x)}{\text{std}(x)}$$

To ensure consistent results, don't use Bessel's correction when implementing this.

6. [Reading] Singular value decomposition is a slow, but highly accurate way of solving least squares problems, particularly ones where $\text{rank}(X)$ is less than the number of columns. In our application, the matrices $\mathcal{V}(x)$ and $\mathcal{H}(x)$ won't be close to rank-deficient, so instead of SVD, QR decomposition could be used for increased performance. The efficient usage of Eigen's [HouseholderQR](#) class is, however, beyond the scope of this assignment.

3 Regression Monte Carlo

[Coding] In this exercise, you will implement the Tsitsiklis–Van Roy [TV01] and the Longstaff–Schwartz [LS01] algorithms to price American options. Further details and discussion of these algorithms can be found in [CLP02] and [Gla13].

3.1 Continuation value

Suppose we're pricing a Bermudan with a discrete set of exercise times $\{0 = t_0, t_1, \dots, t_M = T\}$. In order to decide at any given point t_i whether or not to exercise, the immediate exercise value must be compared to the estimated continuation value $C_i(S)$.

We're going to be using polynomial regression methods implemented in Problem 2, so $C_i(S)$ will be of the form:

$$C_i(S) = \beta_{i0} + \beta_{i1}S + \dots + \beta_{id}S^d \quad (4)$$

To store the coefficients β in an orderly manner, you have been provided with the class [PolynomialMCRegression](#) in `mc_regression.cpp`.

Implement the missing methods in the similar class [HermiteMCRegression](#).

3.2 Backward pricers

We describe how to implement `regression_pricer_backward` as declared in `american_pricers.h`.

```
double regression_pricer_backward(const double spot,
    const arr2& paths,
    const int w, const double strike, const double maturity,
    const double interest_rate, MCRegression& mc_regression,
    const MonteCarloRegressionMethod& method);
```

Let's start by taking a look at some of the arguments.

`method`: Assume that this is set to `Tsitsiklis-VanRoy` for now, we will discuss the `Longstaff-Schwartz` case later.

`w`: The payoff of the option will be $\max(w(S - K), 0)$, so $w = 1$ means call, and $w = -1$ means put.

`mc_regression`: An instance of `PolynomialMCRegression` or `HermiteMCRegression`. This object will estimate the continuation values $C_i(S)$ and store the coefficients β_{ij} seen in equation (4).

`paths`: An array with N rows and M columns. Each of the N rows of `paths` represents a randomly generated path. More formally, its entries are:

$$\begin{aligned} \text{paths}_{k,j-1} &= S_{j\delta t,k} \\ &= S(0) \exp \left\{ (r - q - \sigma^2/2)j\delta t + \sigma\sqrt{\delta t} \sum_{i=1}^j \text{noise}_{Mk+i} \right\} \end{aligned} \quad (5)$$

for $j = 1 \dots M$ and $k = 0 \dots N - 1$, and $\delta t = T/M$. The noise array above will contain an array with NM i.i.d. standard normal samples.

It is recommended to use a “cumulative sum” function to implement this.

In the body of `regression_pricer_backward`, we will need the following arrays, each of length N :

- E will be the immediate **E**xercise value
- C will be the estimated **C**ontinuation value, computed at each step by regression
- P will be the actual **P**ayoff value

The Tsitsiklis–Van Roy algorithm can be summarised by three equations:

$$E_t = \max(w(S_t - K), 0) \quad (6)$$

$$C_t = \text{proj}(P_{t+\delta t} e^{-r\delta t} | \text{span}(f_1(S_t), f_2(S_t), \dots, f_B(S_t))) \quad (7)$$

$$P_t = \begin{cases} E_t & \text{if } t = T \\ \max(E_t, C_t) & \text{if } 0 < t < T \end{cases} \quad (8)$$

The `proj` operator in equation 7 means “calculate the fitted values of a linear regression”. To compute this, please use the `fit_predict` method of the `mc_regression` argument.

When $t = 0$, the stock price is known: $S(0)$. In this case, the orthogonal projection makes C_0 the average of $P_{\delta t}e^{-r\delta t}$ across paths. In this case the computation should be done in the `fit_predict_at_0` method.

The algorithm returns a single value, $P_0 := \max(E_0, C_0)$.

Implement `regression_pricer_backward` with a `for` loop, using the recursive formulation above.

With a small code change, your function can also compute the Longstaff–Schwartz value. The only difference is equation (8) which should be modified to:

$$P_t = \begin{cases} E_t & \text{if } t = T \\ E_t & \text{if } E_t \geq C_t \text{ and } 0 < t < T \\ P_{t+\delta t}e^{-r\delta t} & \text{if } E_t < C_t \text{ and } 0 < t < T \end{cases}$$

Use the `regression_pricer_backward` function's `method` argument to enable the caller to switch between the two algorithms.

3.3 Forward pricer

When running `regression_pricer_backward`, it not only prices the option, but also modifies the supplied `MCRegression` object². After pricing, the regression object contains all the fitted coefficients, so all the information needed to calculate the continuation values $C_i(S)$.

This allows us to easily implement the “forward version” of our pricer:

```
double regression_pricer_forward(const double spot,
    const arr2& paths,
    const int w, const double strike, const double maturity,
    const double interest_rate,
    const MCRegression& mc_regression);
```

As before we will use an array P of length N for the actual payoff values.

For each path k , this method starts at time $t_0 = 0$, and step by step compares the immediate exercise value $K - S$ to the estimated continuation value $C_i(S)$. If at time $j\delta t$ we find that $K - S_{j\delta t,k}$ exceeds $C_j(S_{j\delta t,k})$, the option is *exercised*, we set $P_k = e^{-j\delta t}(K - S_{j\delta t,k})$, break the forward loop and move on to the next path, starting at time 0.

Return the average payoff value across paths: $\text{mean}(P)$.

Your implementation should use both `MCRegression::predict` and `MCRegression::predict_at_0`.

3.4 Calculations

In the attached spreadsheet, you will be asked to price an option using the methods `regression_pricer_backward` and `regression_pricer_forward` with various parameters. Consider a six-month monthly Bermudan put with spot 40.5, strike 44, volatility 20%. Assume that the risk free rate is 4%, dividend rate is 0%.

We will run $N_e = 100$ experiments, each with 10,000 paths.

²As could be seen from the lack of `const` in front of that argument

For the first experiment, generate 10,000 paths with 6 timesteps each using the LCG + inverse transform method with seed $x_0 = 100$. Refer to equation (5) for details.

For each following experiment, generate another 10,000 paths using the last seed from the previous experiment. If you started with $x_0 = 100$, the second seed should be

$$x'_0 := x_{60,000} = 1185163621$$

Record the prices $V_i = 1 \dots N_e$ for each experiment and report³

$$\hat{V} = \frac{1}{N_e} \sum_{i=1}^{N_e} V_i$$

$$\text{s.e.}(\hat{V}) = \frac{1}{\sqrt{N_e}} \sqrt{\frac{1}{N_e} \sum_{i=1}^{N_e} (V_i - \hat{V})^2}$$

For the “out-of-sample” paths start the first experiment at $x_0 = 200$ instead.

Compare your results with those of the binomial tree method with 6,000 timesteps: 4.079018801027898. Report only the differences `regression.price - binomial.tree.price`.

As a sanity check, make sure that for any given Longstaff-Schwartz setup, the in-sample backward and forward pricers produce the same output. Understand why this should be the case and why it's not true for Tsitsiklis–Van Roy.

3.5 Discussion

1. [Writing] Having filled in the spreadsheet, answer the following question about your results:
 - (a) Which one tends to be more accurate: in-sample or out-sample pricing? Is there a bias one way or the other between the two? Give an intuitive explanation for the bias.
 - (b) Do out-sample results have a bias one way or the other (compared to the true value)? Explain where this bias comes from and why it is less noticeable in-sample.
 - (c) Which of Tsitsiklis–Van Roy and Longstaff–Schwartz produce more accurate results in-sample? Is there a bias one way or the other between the two? Find an intuitive explanation for the bias.
 - (d) Does a polynomial regression with the monomial basis benefit from preconditioning? Explain your answer by looking at the difference between backward and forward prices in-sample.
 - (e) Does a polynomial regression with the Hermite basis benefit from preconditioning? Does it benefit from standardization? Or both?

³The standard errors help us understand how much variation to expect from running the pricer with different seeds. They do not feature in pricing.

- (f) Explain why the pricers typically produce higher values as you increase the degree of the polynomial.
 - (g) Which configuration would you use (basis, degree, preconditioning, TVR/LS, forward/backward, in-sample, out-of-sample) in production and why?
2. [Writing] Read the Section 1 and the beginning of Section 2 of Longstaff and Schwartz’s original paper [LS01].
 They use a non-polynomial basis for their regressions. Is this basis orthogonal under any probability distribution? If so, name the distribution. Does their simulation data come from such a distribution?
 Would their choice of basis would be appropriate for a pricing call option?
 3. [Writing] Let’s assume you’ve generated paths and priced an American option.
 How would you calculate a confidence interval that contains the option’s true value with 95% probability. (Without repeatedly running the pricer)
 How would you calculate the option’s delta?

References

- [CLP02] Emmanuelle Clément, Damien Lamberton, and Philip Protter. *An analysis of a least squares regression method for American option pricing*. Berlin, 2002. URL: <https://hdl.handle.net/1813/9176>.
- [Gla13] Paul Glasserman. *Monte Carlo methods in financial engineering*. Vol. 53. Springer Science & Business Media, 2013.
- [LS01] Francis A Longstaff and Eduardo S Schwartz. “Valuing American options by simulation: a simple least-squares approach”. In: *The review of financial studies* 14.1 (2001), pp. 113–147.
- [TB97] Lloyd N. Trefethen and David Bau. *Numerical linear algebra*. SIAM, 1997. ISBN: 978-0-89871-361-9. DOI: 10.1137/1.9780898719574.
- [TV01] J. N. Tsitsiklis and B. Van Roy. “Regression methods for pricing complex American-style options.” In: *IEEE transactions on neural networks* 12 (4 2001), pp. 694–703. ISSN: 1045-9227. DOI: 10.1109/72.935083. ppublish.